

# Predicting Fantasy Points Project Report

Jarett Smith, Max Thompson, Charlie Deaton, and Balin Allred

2023-12-12

## Contents

<b>Executive Summary</b>	<b>2</b>
Results: . . . . .	2
Motivation: . . . . .	2
Value-Add: . . . . .	2
Conclusion: . . . . .	2
<b>Models</b>	<b>3</b>
Data Setup . . . . .	3
Train, Test, CV . . . . .	3
GLMNet Model . . . . .	3
<b>Appendix:</b>	<b>5</b>
Libraries Used . . . . .	5
Code to Produce GLMNet Model . . . . .	5

# Executive Summary

## Results:

Our model was able to better predict player performance compared to an industry standard, ESPN. We chose X model that beat ESPN's RMSE by XX. This represents a XX% beat of ESPN's model.

## Motivation:

Fantasy sports allows the general public to essentially manage their own virtual franchise. Just like real-world sport franchises, fantasy managers draft their own players according to their belief in a player's ability to attempt to beat other teams in an imaginary "league". According to CNN, roughly 60 million people in North America play fantasy sports. Of the various types of sports, fantasy football is the most popular, with roughly 40 million users. With this large amount of participation, the fantasy sports industry as a whole is worth about \$20.3 billion and is expected to continue to grow at 14% annually. This monetary value stems from increased viewership of sporting events and the legalization of sports betting in several countries.

## Value-Add:

With this fairly new industry rapidly growing, our goal is to gain a competitive advantage using data mining methods to better predict player performance. By outperforming industry giants like ESPN, there is real opportunity to fiscally benefit for those who are monetarily invested and opportunity to appeal to recreational participants. In short, this model is valuable to both financially invested individuals and to your "just-for-fun", casual participant.

## Conclusion:

In summary, our model outperforms ESPN estimates of player performance by XX%. In doing so, we are positioned to provide services to a large audience as higher accuracy predictions have real monetary implications to both die-hard and recreational consumers.

Model	RMSE	RMSE St. Dev.
GLMNet	3.181	0.0235

# Models

## Data Setup

```
master <- filter(master, season >= 2014)
master$pick <- cut(master$pick,breaks = seq(0,400,by=32))
master$pick <- ifelse(master$pick == "(288,320]", "Undrafted",master$pick)

current_season$pick <- cut(current_season$pick,breaks = seq(0,400,by=32))
current_season$pick <- ifelse(current_season$pick == "(288,320]", "Undrafted",current_season$pick)
```

## Train, Test, CV

```
set.seed(2023)
train_rows <- sample(nrow(master),round(nrow(master)*0.7,1),replace = FALSE)

train <- master[train_rows,]
test <- master[-train_rows,]

ctrl <- trainControl(method = "cv", number = 5)
```

This sets up train and holdout (or test) data sets to train models on. By using train and holdout data sets, we can ensure that our model is not over-fitting to the training data and examine how well the model generalizes to new data.

## GLMNet Model

### What GLMNet models are and how they work:

GLMNet models use regularized regression to control the complexity of the model and prevent over fitting. These models are particularly useful for high-dimensional data sets where the number of predictors is large relative to the number of observations. The models use traditional linear regression in addition to penalty terms, alpha and lambda, to perform variable selection and handle correlated predictors. The tuning parameter, lambda, controls the strength of regularization, with larger lambda values leading to simpler models. Similarly, larger alpha levels emphasizes lasso regression which also leads to simpler models.

Optimal alpha and lambda levels were found using 5-fold cross-validation. We found the optimal tuning parameters shown below. Essentially, these are the tuning parameters that minimized the RMSE of the model. Given these parameters in our model, the RMSE on the training data set, testing data set, and for 2023 data are shown below:

```
GLM$bestTune
```

```
##      alpha lambda
## 77      1      0
```

```

# Train RMSE
TrainResults<- round(postResample(predict(GLM, newdata = train), train$points_per_game)[1], 3)
# Test RMSE
TestResults<- round(postResample(predict(GLM, newdata = test), test$points_per_game)[1], 3)
# 2023 RMSE
Y2023Results<- round(postResample(predict(GLM, newdata = current_season)
, current_season$points_per_game)[1], 3)

```

	RMSE
TrainResults	3.101
TestResults	3.233
Y2023Results	3.406

Now that we had established the optimal parameters, we would re-train the model on all available data except for 2023 data. Using the optimal tuning parameters and training on the full data set, we saw improved performance in both the test data set and the 2023 data set.

```

# Train RMSE
TrainResults<- round(postResample(predict(GLM_FULL, newdata = train), train$points_per_game)[1], 3)
# Test RMSE
TestResults<- round(postResample(predict(GLM_FULL, newdata = test), test$points_per_game)[1], 3)
# 2023 RMSE
Y2023Results<- round(postResample(predict(GLM_FULL, newdata = current_season)
, current_season$points_per_game)[1], 3)

```

	RMSE
TrainResults	3.114
TestResults	3.151
Y2023Results	3.366

## Appendix:

### Libraries Used

```
library(dplyr)
library(xgboost)
library(DALEX)
library(glmnet)
library(caret)
library(gbm)
library(tidyverse)
library(nflverse)
library(zoo)
library(plotly)
library(visdat)
library(groupdata2)
library(ggplot2)
library(knitr)
```

### Code to Produce GLMNet Model

This code allowed us to choose the best levels for alpha and lambda, which are the two parameters for regularized regression:

```
alpha <- seq(0,1,by=0.1)
lambda <- seq(-3,0,by=0.5)

set.seed(2023); GLM <- train(points_per_game ~ points_pg_ly + targets_pg_ly +
  wopr_pg_ly + pick + air_yards_pg_ly + total_games_ly +
  pass_attempt_difference + total_positional_investment +
  target_dropoff + epa_pg_ly + years_pro + targets_added_this_year +
  is_on_new_team + points_per_snap_ly + targets_per_snap_ly +
  wopr_per_snap_ly + total_snaps_ly + total_passing_tds_ly +
  total_passing_yds_ly + total_passing_fp_ly + is_returning_coach +
  hc_years_with_team + catch_rate_ly + position + fp_dropoff +
  starter_epa_passing_ly + starter_epa_persnap_passing_ly +
  combine_cluster + qbr_ly_bin,
  data=train,
  method="glmnet",
  tuneGrid=expand.grid(alpha = alpha, lambda = lambda),
  metric="RMSE",
  verbose=FALSE,
  trControl=ctrl)
```

After identifying the ideal levels for alpha and lambda via cross-validation, we re-tuned the model to find the ideal coefficients. The new model seen below was tested on all historical data, as well as 2023 data.

```
set.seed(2023); GLM_FULL <- train(points_per_game ~ points_pg_ly + targets_pg_ly +
  wopr_pg_ly + pick + air_yards_pg_ly + total_games_ly +
  pass_attempt_difference + total_positional_investment +
  target_dropoff + epa_pg_ly + years_pro + targets_added_this_year +
```

```

is_on_new_team + points_per_snap_ly + targets_per_snap_ly +
wopr_per_snap_ly + total_snaps_ly + total_passing_tds_ly +
total_passing_yds_ly + total_passing_fp_ly + is_returning_coach +
hc_years_with_team + catch_rate_ly + position + fp_dropoff +
starter_epa_passing_ly + starter_epa_persnap_passing_ly +
combine_cluster + qbr_ly_bin,
data=master,
method="glmnet",
tuneGrid=expand.grid(GLM$bestTune),
metric="RMSE",
verbose=FALSE,
trControl=ctrl)

```